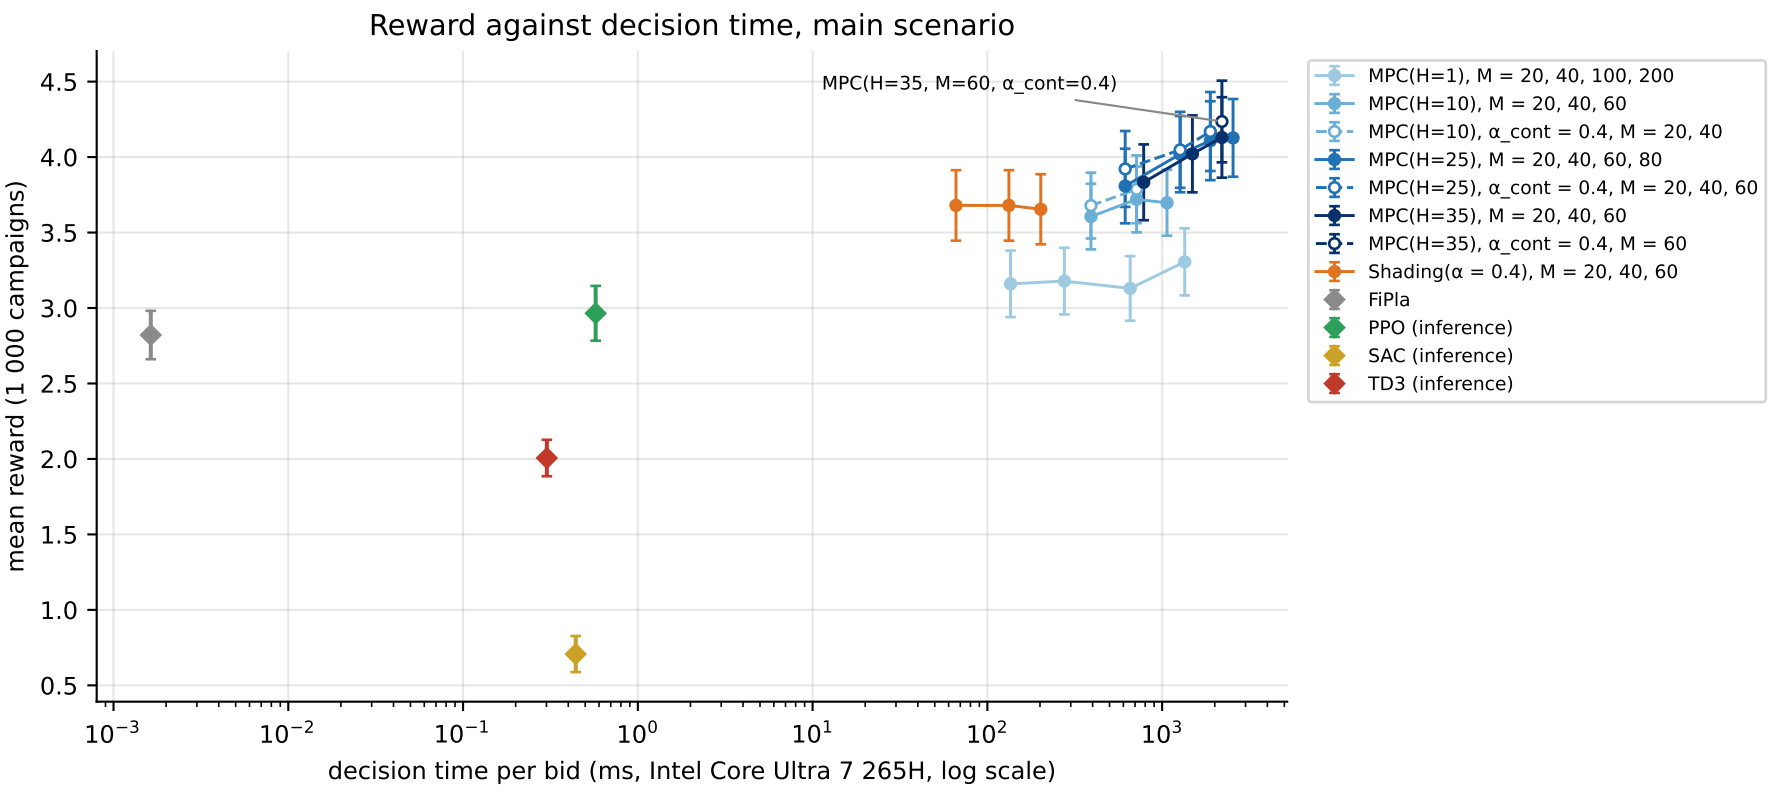
Reward against decision time, main scenario
MPC(H=35, M=60, α_cont=0.4)
mean reward (1 000 campaigns)
decision time per bid (ms, Intel Core Ultra 7 265H, log scale)
MPC(H=1), M = 20, 40, 100, 200
MPC(H=10), M = 20, 40, 60
MPC(H=10), α_cont = 0.4, M = 20, 40
MPC(H=25), M = 20, 40, 60, 80
MPC(H=25), α_cont = 0.4, M = 20, 40, 60
MPC(H=35), M = 20, 40, 60
MPC(H=35), α_cont = 0.4, M = 60
Shading(α = 0.4), M = 20, 40, 60
FiPla
PPO (inference)
SAC (inference)
TD3 (inference)