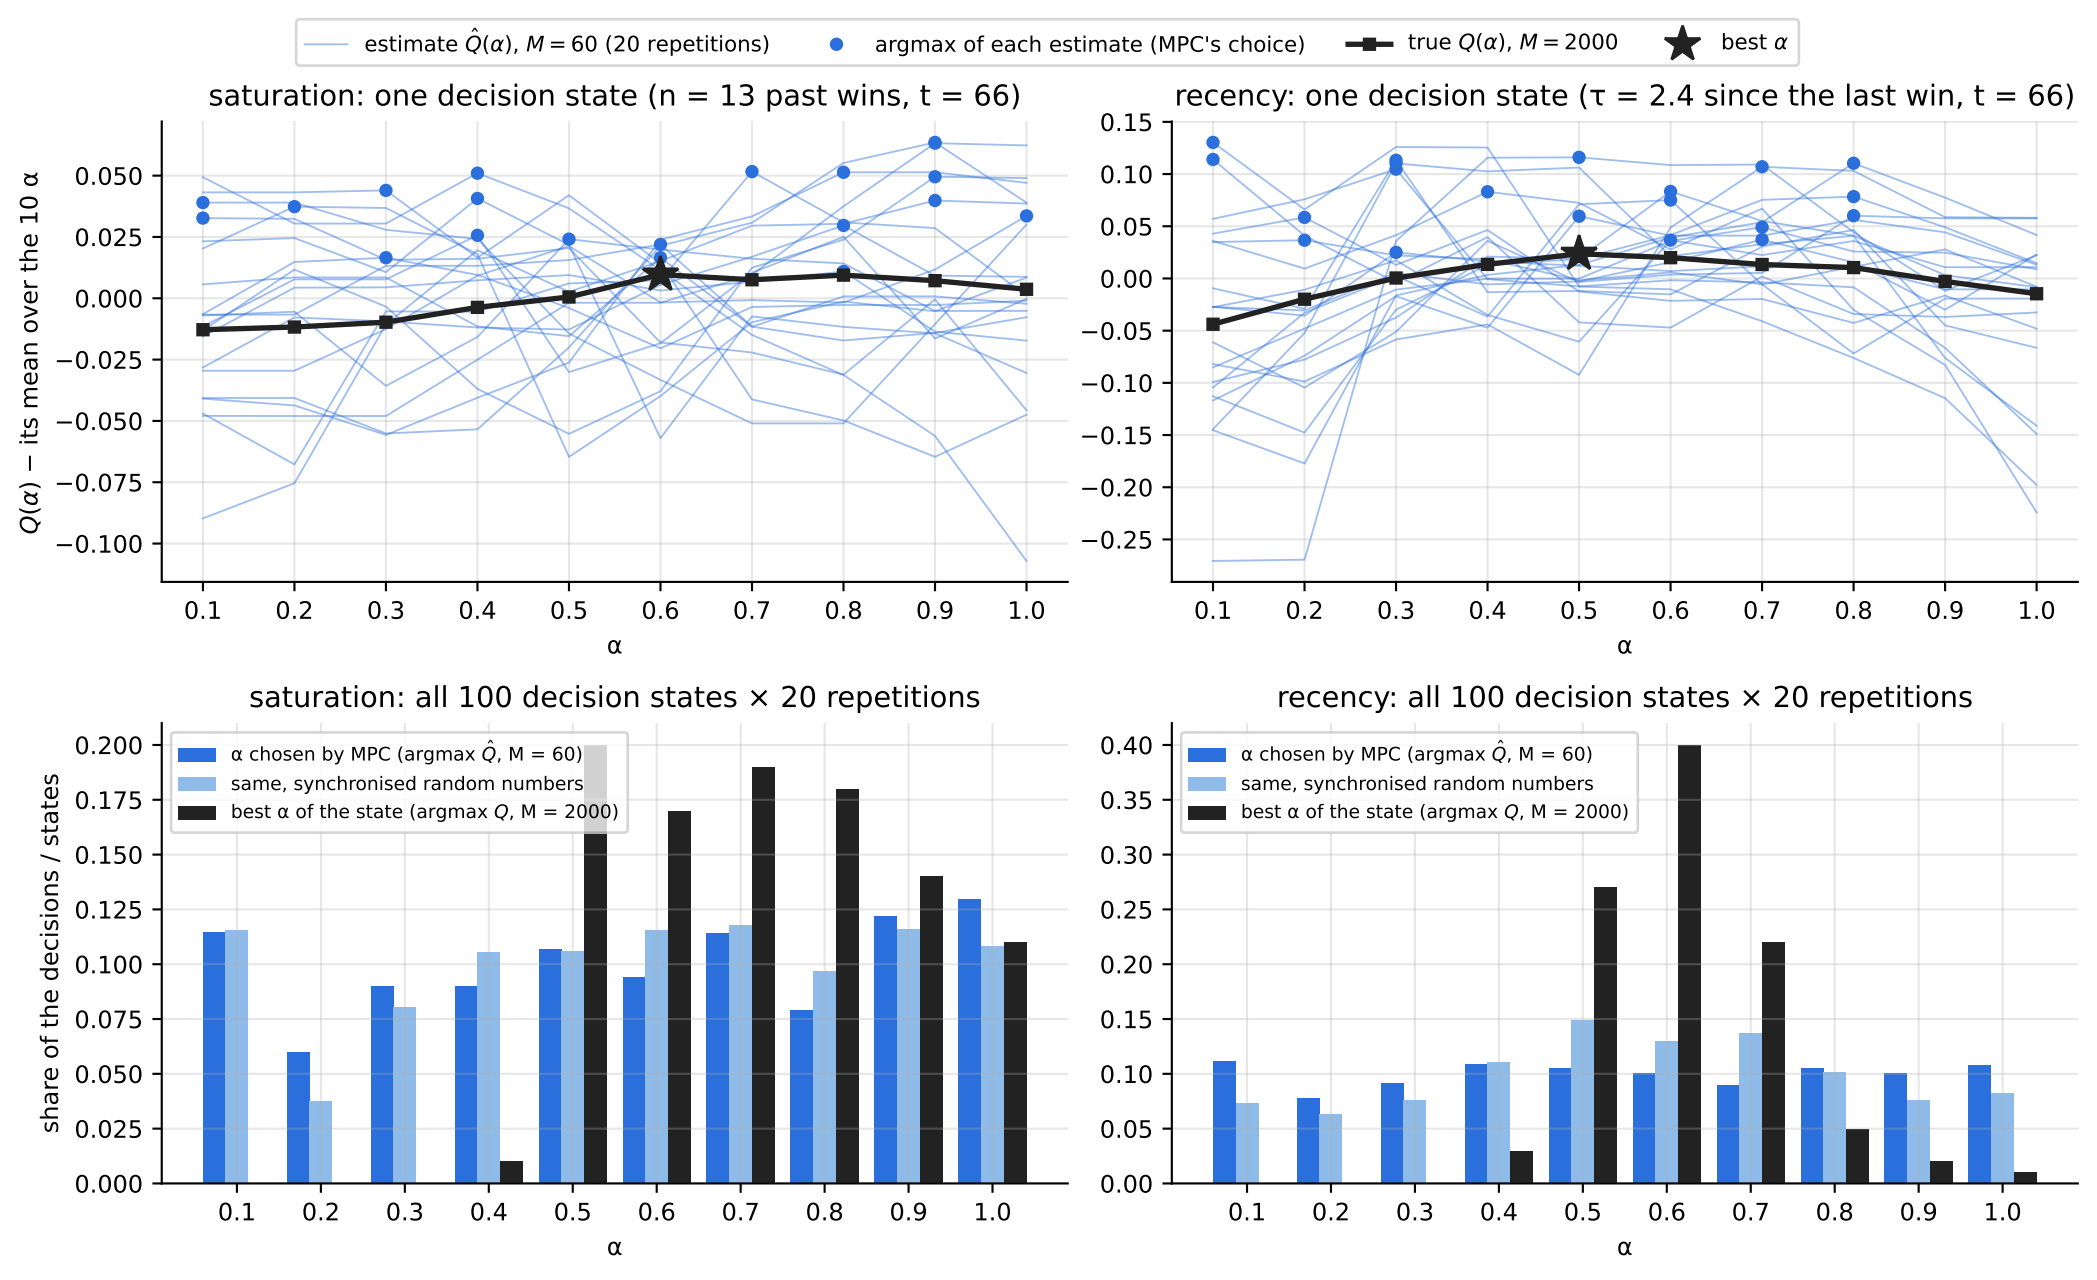
estimate $\hat{Q}(\alpha)$, $M = 60$ (20 repetitions)
argmax of each estimate (MPC's choice)
true $Q(\alpha)$, $M = 2000$
best α
saturation: one decision state (n = 13 past wins, t = 66)
recency: one decision state (τ = 2.4 since the last win, t = 66)
$Q(\alpha)$ − its mean over the 10 α
α
saturation: all 100 decision states × 20 repetitions
recency: all 100 decision states × 20 repetitions
share of the decisions / states
α chosen by MPC (argmax $\hat{Q}$, M = 60)
same, synchronised random numbers
best α of the state (argmax Q, M = 2000)