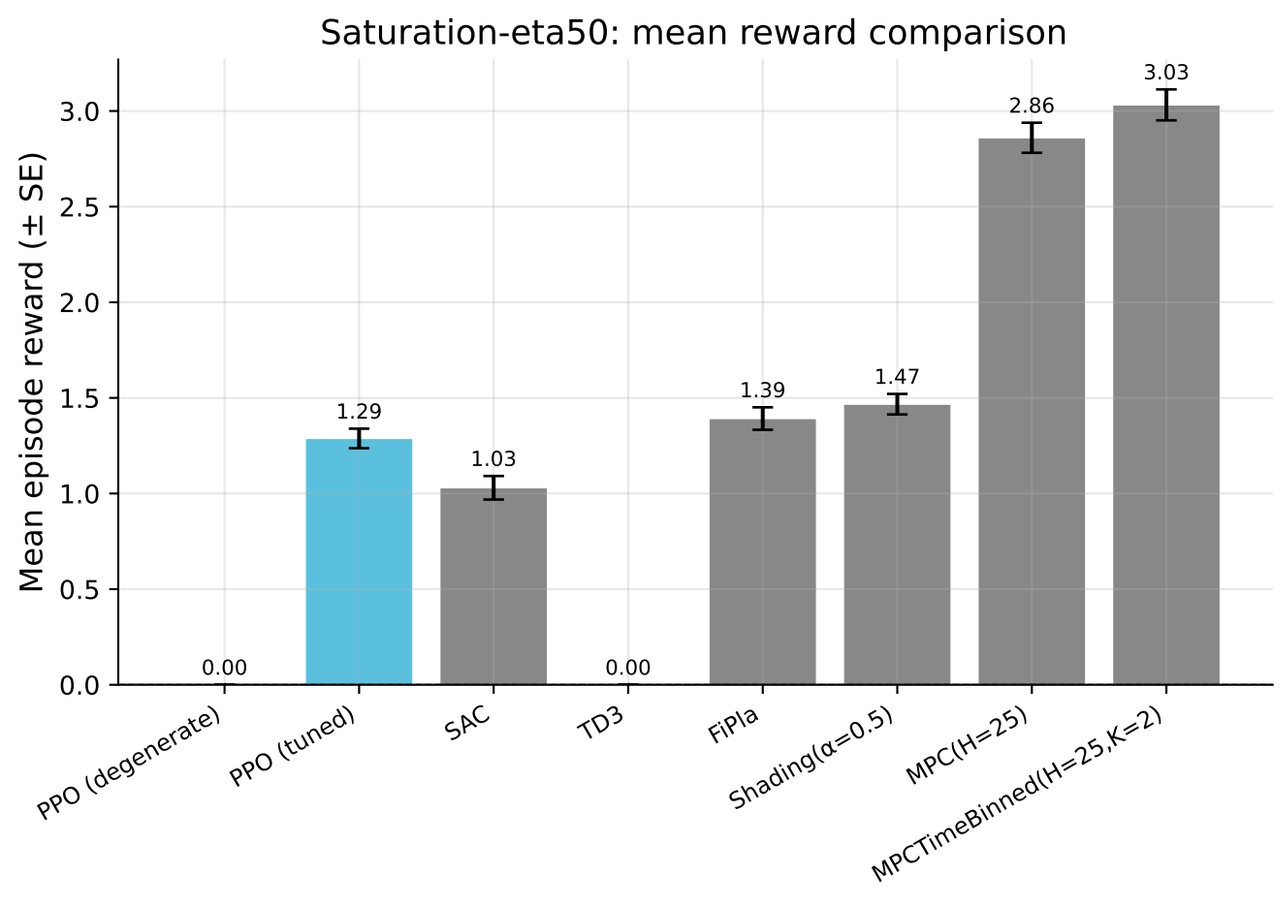
Saturation-eta50: mean reward comparison
Mean episode reward (± SE)
0.0
0.5
1.0
1.5
2.0
2.5
3.0
0.00
1.29
1.03
0.00
1.39
1.47
2.86
3.03
PPO (degenerate)
PPO (tuned)
SAC
TD3
FiPla
Shading(α=0.5)
MPC(H=25)
MPCTimeBinned(H=25,K=2)

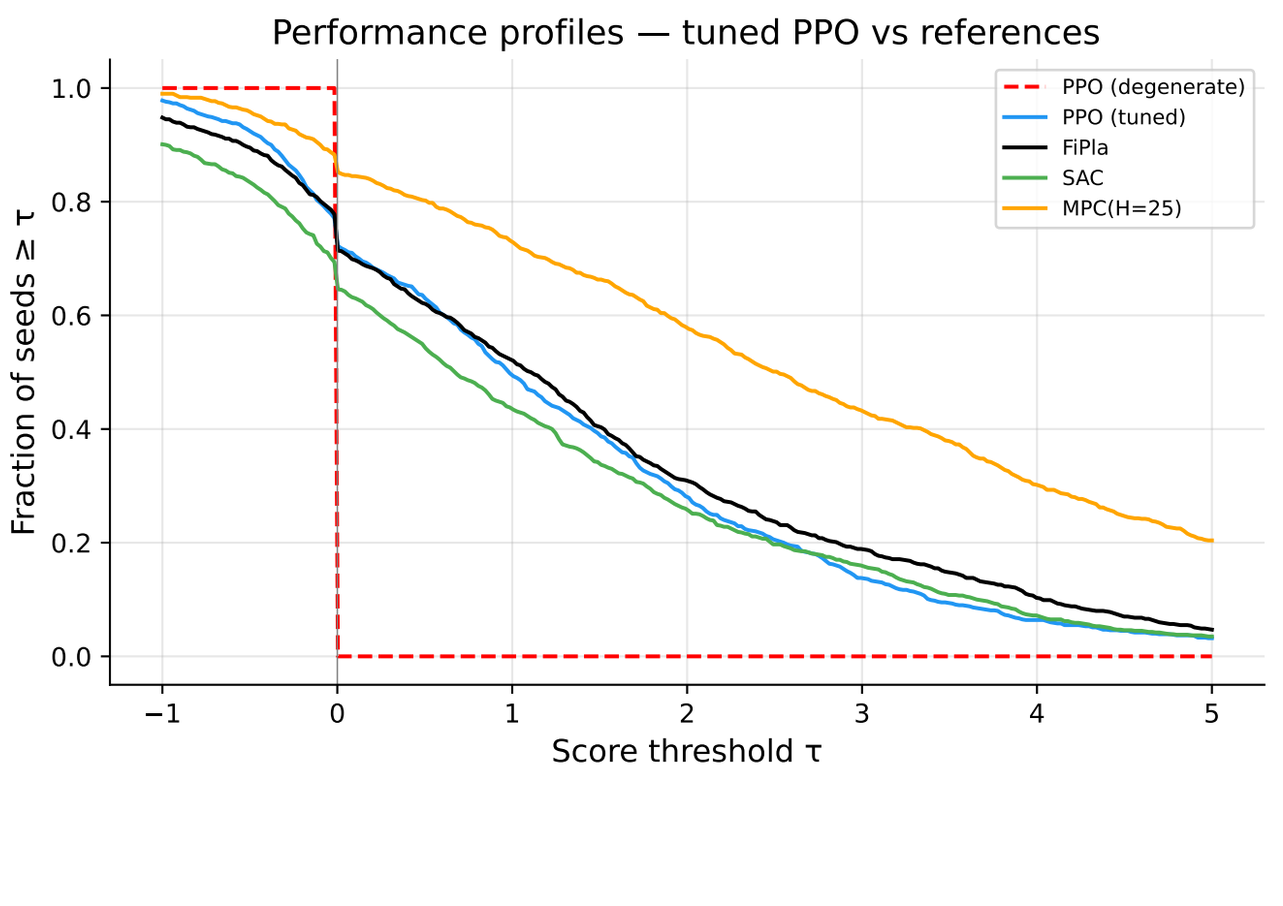
Performance profiles — tuned PPO vs references
Fraction of seeds ≥ τ
0.0
0.2
0.4
0.6
0.8
1.0
−1
0
1
2
3
4
5
Score threshold τ
PPO (degenerate)
PPO (tuned)
FiPla
SAC
MPC(H=25)